

BENV0091 Lecture 8: APIs

Patrick de Mars



Lecture Overview

- String formatting
- APIs:
 - Introduction
 - National Grid Carbon Intensity
 - Met Office

String Formatting

Using paste in a figure title

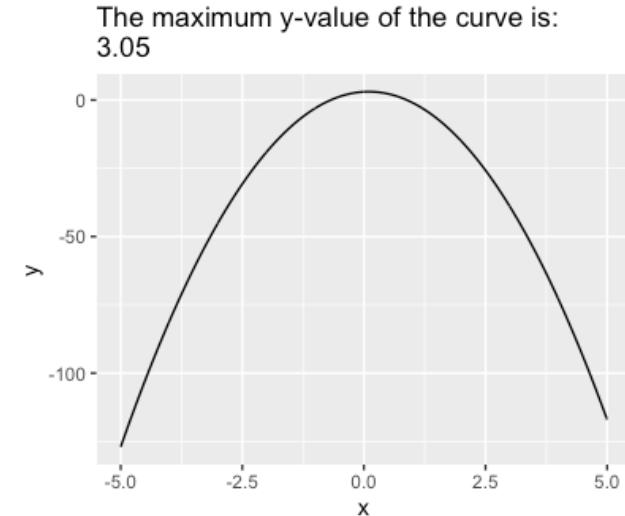
Paste Functions

- The functions `paste()` and `paste0()` are used to join character vectors (strings)
 - `paste()`: arguments are separated by spaces (by default)
 - `paste0()`: arguments are not separated
- Task: use `paste()` to print:
 - "The weather outside is {adjective}"
- Task: use `paste0()` to print:
 - "The time is: {time}"

```
x <- seq(-5, 5, 0.01)
y <- -5*x**2 + x + 3
df <- tibble(x = x, y = y)

title <- paste0("The maximum y-value of the curve is: \n", max(y))

ggplot(df, aes(x = x, y = y)) +
  geom_line() +
  labs(title = title)
```



*Use `now()` to return
the current time*

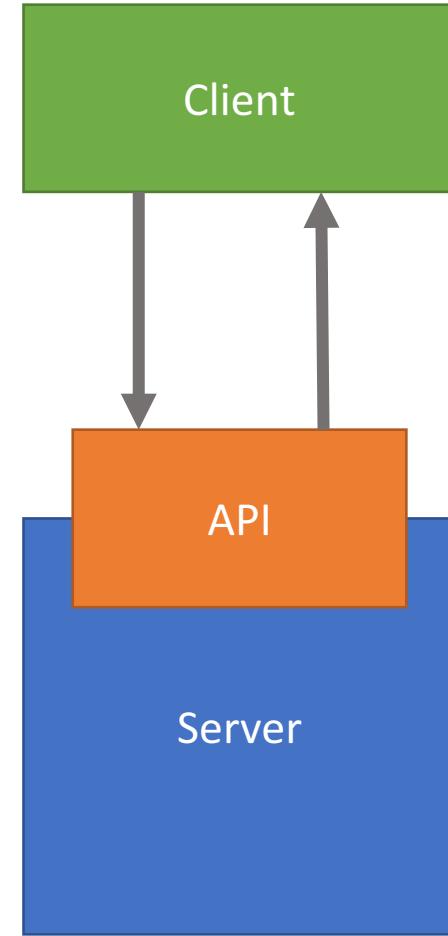
C-Style String Formatting

- It can be challenging or verbose to use paste functions to insert variables in the middle of a string:
 - “Your train departs at HH:MM and is expected to arrive at HH:MM”
 - “The wind speed is {x}mph”
 - The `sprintf()` function enables variables to be included in strings using C-style formatting
 - Task: create a vector of normally distributed values using `rnorm(1000)` then print:
 - “The mean of the data is {mean} and the standard deviation is {sd}”
- ```
sprintf("The wind speed is: %smph", wind_speed)
```
- ```
vals <- rnorm(1000)
```

APIs

APIs

- An **Application Programming Interface** (API) is a set of rules or conventions determining how a user (**client**) can request **resources** from a **server**
- **Server**: runs the API and runs any server-side code
- **Client**: makes **requests** and exchanges data with the server via the API
- **API**: interprets requests to facilitate the exchange of data between client and server



API Examples

- Weather forecasts
- Flight times
- Social media updates

APIs in R

- Hypertext Transfer Protocol (HTTP) is a standard for transferring resources (data, results, code, webpages...) via the Web
- In R, the `httr` package is used for handling HTTP requests and responses
- There are a number of HTTP methods that can be used to specify requests, such as:
 - GET: to request a resource
 - POST: to create a resource
 - PUT: to change a resource
 - DELETE: to remove a resource
- In this lecture we will focus on GET
- Task: install the `httr` package

National Grid's Carbon Intensity API

- The National Grid's Carbon Intensity API is an open-access API that can be used to retrieve forecast and estimated Carbon Intensities for the GB power system
- Task: use GET to make a **request** for the current carbon intensity, assigning the **response** to a variable `r`
- Each API is unique and requires its own methods for extracting information from responses

```
r <- GET('https://api.carbonintensity.org.uk/intensity')
```

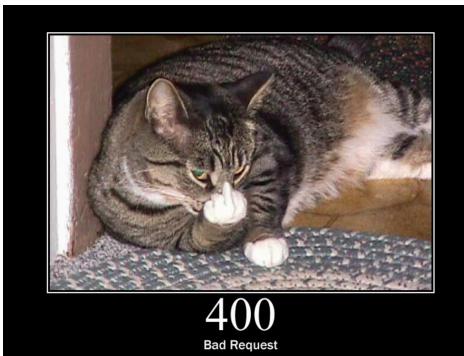
Requesting current carbon intensity

```
> names(r)
[1] "url"          "status_code"   "headers"      "all_headers"  "cookies"      "content"
[7] "date"         "times"        "request"     "handle"
```

Attributes of the response

API Response

- The HTTP response includes the following elements:
 - Status code
 - Headers
 - Content (body)
- The headers contain metadata about the response such as the response format and the file size
- Task: check the **status code** of your response with `r\$status_code`
 - A successful request returns a status code of 200
 - 404 indicates a bad URL



API Response: Body

- The body of an API response contains the information we have requested
- In this case the response is in **JSON format**
- The `content()` function can be used to retrieve the body of a response as **plain text** (or in a raw format)
- However, the text is still not readable in this form: the `prettify()` function (`jsonlite`) can be used to render the plain JSON text to a readable format
- This is readable but still not a convenient R object to handle
- Task: use `prettify()` to print the JSON response (see right)

```
> content(r, 'text', encoding = 'UTF-8')
[1] "{ \r\n  \"data\": [{ \r\n    \"from\": \"2021-11-29T19:00Z\", \r\n    \"to\": \"2021-11-29T19:30Z\", \r\n    \"intensity\": { \r\n      \"forecast\": 194, \r\n      \"actual\": 204, \r\n      \"index\": \"moderate\" \r\n    } \r\n  }] \r\n}"
```

`prettify(content(r, 'text'))`

```
{\n  \"data\": [\n    {\n      \"from\": \"2021-11-02T15:00Z\", \n      \"to\": \"2021-11-02T15:30Z\", \n      \"intensity\": { \n        \"forecast\": 248, \n        \"actual\": 262, \n        \"index\": \"high\" \n      } \n    } \n  ]\n}
```

JSON

- JSON stands for JavaScript Object Notation and is a file format, consisting of key-value pairs
- JSON supports **nested data** such as in the example on the right
 - Nested data cannot be represented in CSV files!
- JSON files are handled with the `jsonlite` package in R
- Task: install `jsonlite`

Example JSON

```
{  
  "team1": {  
    "Rafael": {  
      "age": 35,  
      "hand": "left"  
    },  
    "Roger": {  
      "age": 40,  
      "hand": "right"  
    }  
  },  
  "team2": {  
    "Novak": {  
      "age": 34,  
      "hand": "right"  
    },  
    "Andy": {  
      "age": 34,  
      "hand": "right"  
    }  
  }  
}
```

Response Formats: JSON, XML, CSV

- APIs can respond with any data format so long as it can be serialised
- Most APIs use XML and/or JSON as the response format for data
- *The BMRS API is a notable exception: it returns XML or CSV*

JSON

- Easy to handle with a JSON parser (e.g. jsonlite in R)
- Not as flexible (e.g. does not support metadata)

XML

- Long history and well-known by developers; supports browser rendering
- Difficult to interpret for non-developers

CSV

- Easy to interpret for data scientists
- Does not support hierarchical data

Convert from JSON to an R Object

- The `toJSON()` and `fromJSON()` functions (**jsonlite**) can be used to convert between JSON format and R objects such as data frames or lists
- Task:
 - use `fromJSON()` to convert the response to an R object (in this case a list)
 - Retrieve `data` from the parsed list
 - Finally, use `unnest()` (tidy) to convert the list of data frames to regular columns
- *As you can see, there are necessarily bespoke solutions to parsing API responses*
- Task: wrap the code above in a function that can be re-used

```
r <- GET('https://api.carbonintensity.org.uk/intensity')
parsed <- fromJSON(content(r, 'text'))
df <- parsed$data %>% unnest(cols = intensity)
```

Nested data frame (in JSON format)

```
[  
 {  
   "from": "2021-11-29T20:30Z",  
   "to": "2021-11-29T21:00Z",  
   "intensity": {  
     "forecast": 188,  
     "actual": 182,  
     "index": "moderate"  
   }  
 }  
 ]
```

Unnested data frame (in JSON format)

```
[  
 {  
   "from": "2021-11-29T20:30Z",  
   "to": "2021-11-29T21:00Z",  
   "forecast": 188,  
   "actual": 182,  
   "index": "moderate"  
 }  
 ]
```

Routing and Endpoints

- Endpoints are resources which are available through the API such as:
 - Creating a plot
 - Returning data for a specific date
- Routes determine which endpoint the client is requesting, and is specified in the URL
- Each route can have multiple endpoints corresponding to different verbs: GET, POST etc.



This endpoint returns the carbon intensity between 2020-01-01 and 2020-01-08

More National Grid Functionality

- The Carbon Intensity API [documentation](#) describes endpoints for:
 - Regional carbon intensity
 - National generation mix
- Tasks:
 - Write a function to request the CO₂ intensity data between two dates and parse it into a data frame
 - Write a function to calculate the **average** carbon intensity for the past week
 - Write a function to calculate the period of **highest** carbon intensity in the past week
 - Which regions have the lowest and highest carbon intensities right now?
 - Get the forecast carbon intensity for the next 24 hours (fw24h endpoints)

Datetimes must be in ISO8601 format YYYY-MM-DDThh:mmZ e.g.
2017-08-25T12:35Z

Remember the `sprintf()` and paste functions for string formatting

The `days()` function (*lubridate*) can be used to add or subtract days from datetimes

API Keys

- The National Grid API can be used without authentication
- Many APIs require an authentication key (or **access token**) to make a request
- Note: be careful not to publish API keys on Github! It is safer to store API keys on your local machine (for instance as **environment variables**) and retrieve them when needed
 - Environment variables should be stored in a special file: ` `.Renviron`
 - See ` `Sys.getenv()` for retrieving environment variables
 - [Useful blog post on storing API keys as environment variables](#)

Met Office API

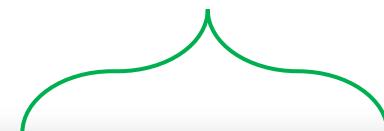
- The Met Office API can be used to retrieve location-specific or national forecasts
 - See [documentation](#)
- You should have already signed up to the Met Office DataPoint – your API Key is available on your [account page](#).
- Task:
 - Use your API key and the URL below to request the site list
 - Parse the response to a data frame
 - Find the location ID for “Land’s End”

```
# Getting site list
url <- sprintf('http://datapoint.metoffice.gov.uk/public/data/val/wxfcs/all/json/sitelist?key=%s', API_KEY)
```

Query Parameters

- So far we have seen **hierarchical parameters** (such as the date or location) separated by ` `/
- However, **query (or optional) parameters** are used by many APIs for routing to endpoints
- Query parameters are separated from the hierarchical parameters by `?`, and separated from *each other* by `&`
- The code below has 1 hierarchical parameter (location ID) and 2 query parameters (resolution and API key)
- Task: use the URL below with your API key and the ID for Land's End to request the 3hourly forecast for Land's End

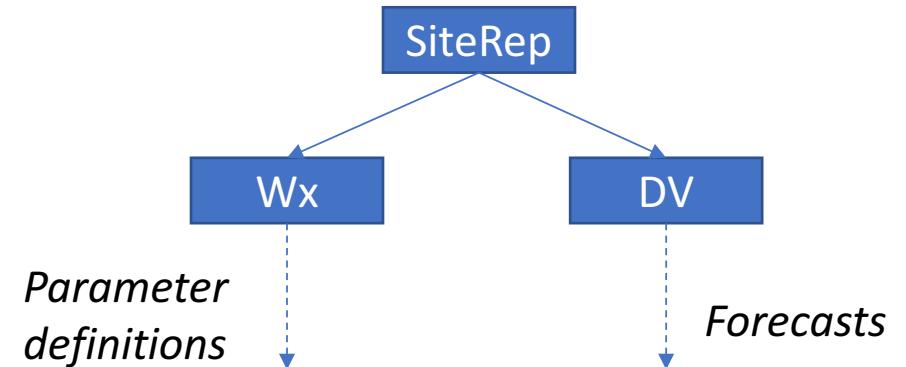
Query parameters



```
url <- sprintf('http://datapoint.metoffice.gov.uk/public/data/val/wxfcs/all/json/%s?res=3hourly&key=%s', location_id, API_KEY)
```

Retrieving the Forecast

- The request in the previous slide returns a 5-day forecast at 3-hourly resolution
- Task: run `prettyf(content(r, 'text'))` to expose the structure of the JSON response
- You should find there are two levels after SiteRep:
 - Wx: defines the parameters
 - DV: contains location data and the forecast



Parameter definitions

```
> data$SiteRep$Wx$Param
```

	name	units	\$
1	F	C	Feels Like Temperature
2	G	mph	Wind Gust
3	H	%	Screen Relative Humidity
4	T	C	Temperature
5	V		Visibility
6	D	compass	Wind Direction
7	S	mph	Wind Speed
8	U		Max UV Index
9	W		Weather Type
10	Pp	%	Precipitation Probability

Retrieving the Forecast

- The data is buried even deeper within the response
- The Rep object is a list of 5 data frames with the forecast for the following 5 days
- Task: use a **for loop** to create a vector giving the Wind Speed for the next 5 days (at 3 hourly intervals)

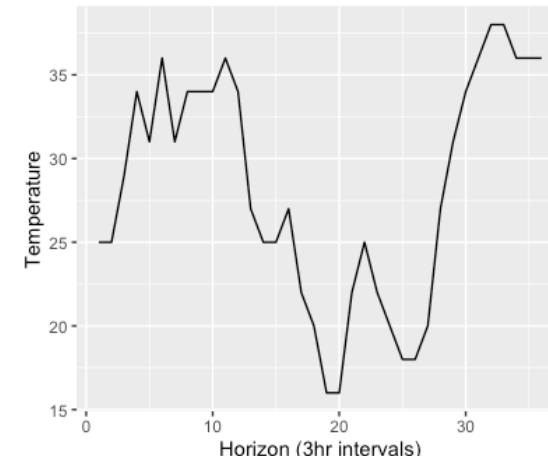
```
> data$SiteRep$DV$Location$Period$Rep
```

```
[[1]]
```

	D	F	G	H	Pp	S	T	V	W	U	\$
1	W	7	34	99	19	25	11	M0	8	1	720
2	W	7	34	99	11	25	11	M0	7	1	900
3	WSW	7	38	91	74	29	11	VG	12	0	1080
4	WSW	7	45	96	98	34	12	G0	15	0	1260

```
[[2]]
```

	D	F	G	H	Pp	S	T	V	W	U	\$
1	WNW	7	40	90	91	31	12	EX	12	0	0
2	NW	6	47	84	36	36	11	EX	7	0	180
3	NW	5	43	75	8	31	10	EX	7	0	360
4	NW	5	43	73	22	34	10	EX	7	1	540
5	NW	4	43	78	43	34	9	EX	10	1	720
6	NW	4	47	71	45	34	9	EX	7	1	900
7	NNW	4	47	72	49	36	9	EX	7	0	1080
8	N	2	45	82	46	34	8	VG	9	0	1260



Period-Specific Forecasts

- By adding an addition query parameter `time`, it is possible to retrieve a forecast for a single period
 - Note: forecasts are only available at 3 hourly intervals: 0h, 3h, 6h,...,21h!
 - Note: the time must be in ISO 8601 format
- Task: obtain the forecast Feels Like Temperature for **tomorrow at midday**
- Task: wrap your code in a function with an argument specifying a time to retrieve the forecast

A lubridate solution to 24 hours ahead in ISO 8601 format

```
time <- (now() + hours(24)) %>% format.ISO8601()
```

```
forecast_tomorrow_feels_like_temp <- function(hour){  
  # Your code goes here...  
  return(feels_like_temp)  
}
```

Some Useful APIs

- Elexon BMRS: [docs \(pdf\)](#)
 - Detailed GB electricity consumption and production data
- Renewables Ninja: [docs](#)
 - Simulated wind and solar PV generation
- Octopus Energy: [docs](#)
 - Tariff charges, electricity/gas smart meter reads etc.
- Open Weather: [docs](#)
 - Current, historical and forecast weather data
- Twitter: [docs](#)
 - Retrieve tweets, create tweets etc.
- European Environment Agency: [docs](#)
 - Environmental indicators, water quality etc.
- UK public sector APIs: [catalogue](#)
 - Companies House, NHS 999 callouts etc.
- OECD APIs: [catalogue](#)
 - Economic indicators

Building an API in R

- The **plumber** package can be used to build your own web API in R (a great way to improve your understanding of APIs!)
- You can test this out by running it **locally** – making it publicly available can be achieved on RStudio Connect (at a cost)

